

ParSplice Keyspace Locality

1 INTRODUCTION

The fine-grained data annotation capabilities provided by key-value storage is a natural match for many types of scientific simulation. Simulations relying on a mesh-based decomposition of a physical region may result in millions or billions of mesh cells. Each cell contains materials, pressures, temperatures and other characteristics that are required to accurately simulate phenomena of interest. In our target application, the ParSplice [1] molecular dynamics simulation, a hierarchy of cache nodes and a single node key-value store are used to store both observed minima across a molecule’s equation of motion (EOM) and the hundreds or thousands of partial trajectories calculated each second during a parallel job. Unfortunately, if we scale the system the IO to the storage hierarchy will quickly saturate both the storage and bandwidth capacity of a single node.

In this paper we present a detailed analysis of how the ParSplice application accesses key-value pairs over the course of a long running simulation across a variety of initial conditions. We reason that limiting the size of a cache on a single node saves memory and sacrifices negligible performance. This type of analysis (1) shows the capacity and resource requirements of a single node and (2) will help inform our load balancing policies for when we switch to a distributed key-value store back-end to store EOM minima. We need to know when and how to partition the keyspace: a smaller cache hurts performance because key-value pairs need to be retrieved from other nodes while a larger cache has higher memory pressure.

To explore the effects of different cache management strategies, we link the Mantle policy engine into ParSplice [1]. Developers write policies for “when” they want data moved, “where” they want data moved, and “how much” of the data to move and the framework executes these policies whenever a decision needs to be made. This abstraction helps developers unfamiliar with the domain quickly reason about, develop, and deploy new policies that control temporal and spatial locality. It has already proven to be a critical control plane for improving file system metadata load balancing [2] and in this work we show its usefulness in cache management for the changing key-value workloads generated by ParSplice.

We show how Mantle:

- decomposes cache management into independent policies that can be dynamically changed, making the problem more manageable and facilitating rapid development. Changing the policy in use is critical in applications such as ParSplice that have alternating stable and chaotic simulation “access regimes” over the course of a long-running simulation.

- has useful primitives that, while designed for file systems, turn out to also be effective for cache management. This finding shows how the Mantle engine generalizes to a different domain and code-base.
- can be used to quickly deploy a variety of cache management strategies, ranging from basic algorithms and heuristics to statistical models and machine learning.

This last contribution is explored in Section §4, where we try a range of policies from different disciplines; but more importantly, in Section §5, we conclude that the collection of policies we designed for ParSplice’s cache management are very similar to the policies implemented in the Ceph file system metadata load balancer. This lays the foundation for future work, where we will focus on formalizing a collection of general data management policies that can be used across domains and services. The value of such a collection eases the burden of policy development and paves the way for solutions that remove the administrator from the development cycle, such as (1) adaptable policies that automatically switch to new strategies when the current strategy behaves poorly (e.g., thrashing, making no progress, etc.), and (2) policy generation, where new policies are constructed automatically by examining the collection of existing policies. Such work is made possible with Mantle’s ability to dynamically change policies

2 PARSPLICE KEYSPEC ANALYSIS

We demonstrate these contributions by changing the input to our molecular dynamics simulation, which changes the keyspace access patterns.

2.1 Structured Access Regimes

Figures 4, 5 and 6.

- change immediately, different sizes
- monotonically increasing
- random access to a single key
- early keys accessed more

Conclusion: unique patterns of a real HPC application

2.2 Cache Size Analysis

Figure 7.

Conclusion: workload phases (high request rate to a small number of keys and then low request rate to a large number of keys) need a dynamic load balancing policy.

2.3 Overfitting Policies

Figures 8 and 9.

Side Idea: we need to figure out what ParSplice memory is sensitive to: max usage or usage over time.

Conclusion: dynamic policies absorb the cost of a high read request rate for a 2.5 hour run, but it is infeasible to do this for every combination of system setup, job lengths, parsplice parameters.

Discussion. Why don't we just an LRU cache:

- we can do better (workload is structured and has locality)
- finding the size of the cache is hard
- hotspots dissipate too quickly

3 MANTLE ENGINE: METHODOLOGY

3.1 Experimental Setup

3.2 Mantle: Dynamic Load Balancing Policies

- motivation: Mochi load balancer microservice
- background: CephFS implementation
- library architecture, callbacks, environment

3.3 Integrating Mantle into ParSplice

- providing environment of metrics
- identifying where policies are made

4 MANTLE BRAINS: TOOLS WE PLUG IN TO DETECT “WHEN”

Requirements: run online, be fast enough to run as often as we want to detect regimes

4.1 System Specific Knowledge

e.g., request rate, unique keys in a sliding window, bandwidth capabilities. For example, we know that LevelDB cannot handle high IO request rates.

4.1.1 Request Rate.

4.1.2 Belady's Min.

4.2 Domain Specific Knowledge

e.g., ParSplice key access locality.

4.2.1 Regime Detection. At each time step, we find the lowest ID and compare against the local minimum, which is the smallest ID we have seen thus far. If we move left to right, the local minimum never changes because the local minimum will start small. If we move from right to left, the local minimum changes at each access regime. For points z , y , and x , if the local minimum is the same we are in a regime. Processing y , we set the local minimum to be $\min(y, m_l)$, where m_l is the local minimum of the previous time step of z . The algorithm incorrectly detects a regime change if the local minimum of y is lower than local minimum of z , since y may have points *within* z ; recall that we are trying to detect the whole fan, not just the bottom edge of each fan.

4.2.2 Trajectory Length.

```

1  -- assume that (ts, keyid) are in a table
2  local function when()
3
4
5      if servers[whoami]["load"] > target then
6          if servers[whoami]["load"] > absorb_target then
7              WRstate(1)
8          end
9          if RDstate() == 1 then
10             return true
11          end
12      end
13      return false
14  end

```

Figure 1: ParSplice cache management policy.

4.3 Failed, Overcomplicated Brains

These techniques proliferated more, less transparent knobs

- Statistics
- Calculus
- K-Means
- DBScan
- Anomaly Detection

4.4 Cloud Techniques: Elastic Search

What if we re-provision resources in response to events outside the application's control, such as a slow Lustre.

4.5 Future work

How Much: cache policy from past, regime detection

5 RELATION TO FILE SYSTEMS

The code snippets in Figures 2 and 3 are the policies used in ParSplice and CephFS, respectively. ParSplice uses policies to manage its caches and CephFS uses policies to control load balancing, but they both can be expressed with the Mantle API. From a high-level the ParSplice policy trims the cache if the cache reaches a certain size *and* if it has already absorbed the initial burstiness of the workload; the CephFS policy migrates load if the metadata load is higher than the average load *and* the current load has been overloaded for more than two iterations.

Condition for “Overloaded” (Fig. 2:Line 2; Fig. 3:Line 2) - these lines detect whether the node is overloaded using the “load” calculated in the load callback; while the load calculations and thresholds are different, the actual logic is exactly the same. Recall that this decision is made locally because there is no global scheduler or centralized intelligence.

State Persisted Across Decisions (Fig. 2:Lines 5,6; Fig 3:Lines 3,4,10) - these lines use the Mantle API to write/read state from previous decisions. For ParSplice, we save a boolean that indicates whether we have absorbed the workload's initial burstiness. For CephFS, we save the number of consecutive instances that the server has been

```

1 local function when()
2   if servers[whoami]["load"] > target then
3     if servers[whoami]["load"] > absorb_target then
4       WRstate(1)
5     end
6     if RDstate() == 1 then
7       return true
8     end
9   end
10  return false
11 end

```

Figure 2: ParSplice cache management policy.

```

1 local function when()
2   if servers[whoami]["load"] > target then
3     overloaded = RDstate() + 1
4     WRstate(overloaded)
5     if overloaded > 2 then
6       return true
7     end
8   end
9   else then
10    WRstate(0)
11  end
12  return false
13 end

```

Figure 3: CephFS file system metadta load balancer.

overloaded. We also clear the count (Line 10) if the server is no longer overloaded. The underlying implementation saves the values to local disk.

Condition that Switches Policy (Fig. 2:Line 3; Fig. 3:Line 5) - these lines switch the policies using information from previous decisions. ParSplice trims its cache once it eclipses the “absorb” threshold while CephFS allows balancing when overloaded for more than two iterations. The persistent state is essential for both of these policy-switching conditions.

- Lustre Trace
- LinkedIn Trace
- Nathan’s Trace

5.1 Using File System Balancers for ParSplice

5.2 Using ParSplice Balancers for File Systems

5.3 Visualizing File System Traces like ParSplice Keyspace Traces

6 CONCLUSION

REFERENCES

- [1] Danny Perez, Ekin D Cubuk, Amos Waterland, Efthimios Kaxiras, and Arthur F Voter. Long-Time Dynamics Through Parallel Trajectory Splicing. *Journal of chemical theory and computation* (????).




Figure 4

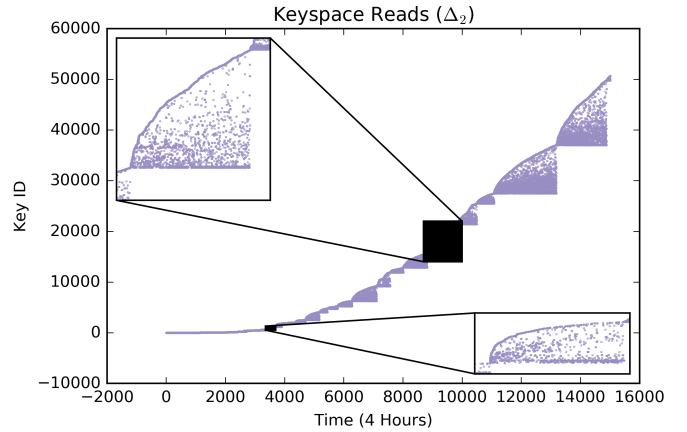


Figure 5

- [2] Michael A. Sevilla, Noah Watkins, Carlos Maltzahn, Ike Nassi, Scott A. Brandt, Sage A. Weil, Greg Farnum, and Sam Fineberg. 2015. Mantle: A Programmable Metadata Load Balancer for the Ceph File System. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '15)*.




Figure 6




Figure 9




Figure 7




Figure 8